

# Linear equations and iterative methods

## Judd Chapter 3

David Childers (thanks to Y. Kryukov, K. Judd, and U.  
Doraszelski)

CMU, Tepper School of Business

March 15, 2023

# Numerical Linear Algebra

- Linear algebra is basic building block of numerical mathematics
- Algorithms extremely well developed: every language has low and medium level library implementing efficient algorithms
  - *BLAS*: Basic Linear Algebra Subprograms to implement basics, eg OpenBLAS/MKL
  - Mid-level: LAPACK almost universal for standard implementations of applied algorithms
  - High-level: Julia [LinearAlgebra](#), Python [scipy/numpy](#) `linalg`
  - Specialized libraries: (eg for deep learning) reimplement to take advantage of hardware (GPU, etc) or extra structure
- You should understand building blocks of these algorithms: speed, robustness, accuracy in different situations
- Core problems
  - $A*B$  /  $A @ B$ ,  $A \setminus b$ ,  $\text{eig}(A)$  /  $\text{eigen}(A)$ , etc.
  - Many implementations each, with different tradeoffs

# Warmup: Matrix Multiplication

- Let  $A, B$  be  $n \times n$  matrices
- Naive algorithm follows from definition of  $A * B$

$$[A * B]_{ij} = \sum_{k=1}^n A_{ik} B_{kj}$$

- Takes  $n$  multiplications and additions for each of  $n^2$  entries
  - $O(2n^3)$  operations total
- Coppersmith-Winograd algorithm is  $\approx O(n^{2.37})$ 
  - But constant large: only use if  $n$  huge
- Default implementations have cost in between (e.g.  $O(n^{2.81})$ )
- Matrix-vector multiply  $A * b$ ,  $b$   $n \times 1$  is  $O(2n^2)$  by same logic

# Exploiting Structure

- Many matrices arising in practice have additional structure
  - Symmetric, Diagonal, Toeplitz, triangular, sparse, low rank, etc
- Sparse matrices have  $m \ll n^2$  nonzero entries
- Toeplitz & circulant matrices occur under invariances
- Fast algorithms exist which depend on structure: e.g.
  - Diagonal (or tridiagonal) matrix multiplication is  $O(n)$
  - Circulant matrix-vector multiplication is  $O(n \log(n))$ : FFT
  - Will use this for interpolation, etc
- Some algorithms work for all matrices, faster when structured
- Others need explicit info about structure: use *types* to represent
- See `scipy.linalg` or [JuliaLinearAlgebra](#) library families

# The Plan: systems of linear equations

- Direct methods – reliable but slow
  - Back substitution
  - LU decomposition
  - QR and Cholesky decompositions
- Measuring precision: Eigenvalues and condition numbers
- Iterative methods:
  - Classic version – Gauss-Jacobi
  - An improvement – Gauss-Seidel
  - Convergence, dampening and acceleration
  - Alternatives
- Technical improvements: parallelization, randomization, sparsity
- Application: Eigenvectors and Markov chains

# Systems of linear equations

$$Ax = b,$$

where  $A$  is an  $n \times n$  matrix,  $x$  and  $b$  are  $n \times 1$  vectors

- Some simple economic problems are linear:
  - General Equilibrium with linear supply and demand
  - OLS estimator:  $(X'X)\hat{\beta} = X'Y$
- Used as a building block in other methods
  - Newton's method for nonlinear equations or optimization
  - Polynomial or spline coefficients in approximation
  - Value function in policy iteration
  - Ergodic (stable) distribution of Markov Chain
- Simple example used to illustrate iterative methods

# Back-substitution

Suppose  $A$  is lower triangular, i.e.,

$$A = \begin{pmatrix} a_{11} & 0 & \cdots & 0 \\ a_{21} & a_{22} & \cdots & 0 \\ \vdots & \vdots & \ddots & \vdots \\ a_{n1} & a_{n2} & \cdots & a_{nn} \end{pmatrix}.$$

Then we can easily solve it:

$$\begin{aligned} x_1 &= \frac{b_1}{a_{11}}, \\ x_k &= \frac{b_k - \sum_{j=1}^{k-1} a_{kj}x_j}{a_{kk}}, \quad k = 2, 3, \dots, n. \end{aligned}$$

# LU decomposition

- Find  $L$  and  $U$  such that:

$$A = LU$$

- $L$  is **L**ower triangular
- $U$  is **U**pper triangular
- Matlab/Julia (LinearAlgebra)/Scipy: `lu(A)`
- Solving system of equations  $Ax = b \Leftrightarrow L[Ux] = b$ :
  - Find  $L$  and  $U$
  - Solve  $Lz = b$  for  $z$  by back-substitution
  - Solve  $Ux = z$  for  $x$  by back-substitution
  - Matlab/Julia: `x = A\b` (backslash or "left division")  
Numpy/scipy `solve(A,b)`



# Other Decompositions

- **QR decomposition:**  $A = QR$ .
  - $Q$  is orthogonal ( $Q'Q$  is an identity matrix),
  - $R$  is upper triangular.
  - Matlab/Julia/Numpy/Scipy: `qr(A)`
- Solution:  $Ax = b \Leftrightarrow Q'QRx = Q'b$ ,
  - $Q'QR$  is upper triangular  $\Rightarrow$  solve for  $x$  by back-substitution
  - One back-substitution, but decomposition takes longer
- **Cholesky decomposition** ("square root" of  $A$ )
  - Exists if  $A$  is symmetric and positive definite
  - $A = CC'$ ,  $C$  is lower triangular ( $\Rightarrow C'$  – upper).
  - Matlab: `chol(A)` Julia/Numpy/Scipy: `cholesky(A)`

# A bit of theory: matrix analysis

- $\lambda \in \mathbb{C}$  is an **eigenvalue** of  $n \times n$  matrix  $A$  iff:
  - there is an eigenvector  $v \in \mathbb{C}^n$ ,  $v \neq 0$ , ...
  - ... such that  $Av = \lambda v$ , and  $\det(A - \lambda I) = 0$ .
- **Spectrum**  $\sigma(A)$  = set of all  $n$  eigenvalues
  - Matlab/Numpy/Scipy: `eig(A)` Julia: `eigvals(A)`
- **Spectral radius**  $\rho(A) := \max |\sigma(A)|$
- **Norm**:  $\|A\| = \max_{x \neq 0} \frac{\|Ax\|}{\|x\|} = \max_{\|x\|=1} \|Ax\| = \text{norm}(A)$   
/ `opnorm(A)`
- **Condition number**:  $\text{cond}(A) = \|A\| \cdot \|A^{-1}\| = \text{cond}(A)$ 
  - If  $A$  is singular, then the condition number is  $+\infty$ .
  - If  $A$  is near-singular, condition number is high ( $> 10^{10}$ )
- **Spectral condition number**:  $\text{cond}_*(A) = \rho(A) / \min |\sigma(A)|$ 
  - $\text{cond}(A) \geq \text{cond}_*(A)$
  - They tend to have same order of magnitude ( $m$  in  $10^m$ )

# Bounding the errors

- Perturb the system by  $r$ :  $A\tilde{x} = b + r$  (e.g. due to rounding)  
 $\implies$  Error  $e = \tilde{x} - x$ .
- Elasticity of solution w.r.t. error:

$$\frac{1}{\text{cond}(A)} \leq \frac{\|e\|/\|x\|}{\|r\|/\|b\|} \leq \text{cond}(A).$$

- Rule of thumb: each order of magnitude in condition number loses one decimal digit of accuracy from  $x$
- Since we only care about orders of magnitude, we can use  $\text{cond}_*(A)$  instead – it is a lot faster to compute.
- $\Rightarrow$  Pre-conditioning: pick  $D$  so  $DA$  is well-conditioned, solve

$$(DA)x = (Db)$$

# Speed of computation

- LU decomposition:
  - Decomposition:  $n^3/3$  multiplications and divisions
  - Backward substitution:  $n^2$  multiplications and divisions
- QR decomposition – also  $Kn^3$  operations
- Cramer's rule (determinants) =  $n!$  operations = very slow
  - It is useful in theory work, as an expression for solution
- We can often improve speed at cost of a bit of precision
  - $\implies$  Iterative methods

# Gauss-Jacobi iterations

- Idea: solve  $i$ -th equation for  $x_i$  alone:

$$\sum_{j=1}^n a_{ij}x_j = b_i \\ \Rightarrow x_i = \frac{1}{a_{ii}} \left\{ b_i - \sum_{j \neq i} a_{ij}x_j \right\}$$

- Use it to iterate: have guess  $x^k$ , compute  $x^{k+1}$  as:

$$x_i^{k+1} = \frac{1}{a_{ii}} \left\{ b_i - \sum_{j \neq i} a_{ij}x_j^k \right\}, \quad i = 1, \dots, n$$

- Need a starting guess  $x^0$ .
- Need a stopping rule, e.g.:

$$\frac{\|x^{k+1} - x^k\|}{\|x^k\| + 1} < \delta$$

# Gauss-Seidel

- Gauss-Jacobi computes the last element ( $x_n^{k+1}$ ) as:


$$x_n^{k+1} = \frac{1}{a_{nn}} \left\{ b_n - \sum_{j=1}^{n-1} a_{nj} x_j^k \right\}$$

- If we update elements in natural order ( $1, \dots, n$ ) we already have  $x_1^{k+1}, \dots, x_{n-1}^{k+1}$  computed.
- It would speed things up to use them rather than  $x^k$
- Gauss-Seidel does the same for every element  $x_i^{k+1}$ :

$$x_i^{k+1} = \frac{1}{a_{ii}} \left\{ b_i - \sum_{j=1}^{i-1} a_{ij} x_j^{k+1} - \sum_{j=i+1}^n a_{ij} x_j^k \right\}$$

- See Figure 3.2 in the textbook for an illustration.

## Figure 3.2



Gauss-Jacobi (ADGH) versus Gauss-Seidel (ABELN..)

# Convergence and operator splitting

- Operator splitting:  $A = (N - P)$

$$\begin{aligned} Nx &= b + Px \\ x^{k+1} &= N^{-1}(b + Px^k) \end{aligned}$$

- $N$  is selected so it is easy to invert:
  - GJ:  $N$  is the diagonal of  $A$
  - GS (using natural order):  $N$  is the lower triangle of  $A$
- GJ/GS are linearly convergent at rate  $\rho(N^{-1}P)$
- Condition for convergence:  $\rho(N^{-1}P) < 1$ 
  - $\iff$  all eigenvalues are inside the unit circle
  - $\Leftarrow$  Matrix  $A$  diagonally dominant:  $|a_{ii}| > \sum_{j \neq i} |a_{ij}|$
  - 1-dimensional interpretation: slope less than 1



# Precision under linear convergence

- Linear convergence (plus some conditions) imply:

$$\|x^{k+1} - x^k\| \leq \beta \|x^k - x^{k-1}\|$$

- where  $\beta$  is the linear convergence rate
- It follows that  $\|x^k - x^*\| \leq \|x^{k+1} - x^k\| / (1 - \beta)$ 
  - To prove, use triangle inequality
- So, if we want to ensure  $\|x^k - x^*\| < \delta$ , stop when

$$\|x^{k+1} - x^k\| \leq \delta (1 - \beta)$$

- Since we do not know  $\beta$ , we can estimate:  
 $\hat{\beta} = \max_{\kappa=1, \dots, k} \{ \|x^\kappa - x^{\kappa-1}\| / \|x^{\kappa-1} - x^{\kappa-2}\| \}$
- Of course, this approach is valid only if there is convergence

# Guarantees for iterative methods

- Gerschgorin circle theorem:
  - Let  $A$  have eigenvalues  $\{\lambda_s\}_{s=1}^n$ ,  $R_i = \sum_{j \neq i} |a_{ij}|$
  - Then  $\lambda_s \in \cup_{i=1}^n B(a_{ii}, R_i) \forall s$
- Corollary: If  $A$  is diagonally dominant with  $\beta = \sup\{|x| : x \in \cup_{i=1}^n B(a_{ii}^{-1}, R_i)\} < 1$ , then Gauss-Jacobi converges to within error  $\delta$  in  $O(n^2 \log(\frac{1}{\delta}))$  operations
- Informal proof:
  - Each iteration is a diagonal matrix multiply, which is  $O(n^2)$
  - Gerschgorin  $\implies \rho(N^{-1}P) \leq \beta < 1$ ,  $\implies$  linear convergence
- Compare  $O(n^3)$  for direct methods
  - Substantial improvement even for machine precision error
- For G-S, can show  $A$  diagonal dominant OR symmetric positive definite  $\implies \rho(N^{-1}P) < 1$ 
  - Occurs in least squares problems, many 2nd order PDEs
- These structures are sufficient, not necessary

# Dampening: a tweak to iterations

- Gauss-Jacobi algorithm takes steps  $\Delta x$  at each iteration:
  - Each iteration computes  $x^{k+1} = N^{-1}Px^k + N^{-1}b = Gx^k + d$
  - Step is  $\Delta x^{k+1} = x^{k+1} - x^k = Gx^k + d - x^k$
- We can choose to scale  $\Delta x^{k+1}$  by  $\omega$ :

$$x^{k+1} = \omega[Gx^k + d] + (1 - \omega)x^k$$

- $\omega < 1$ : Stabilization / dampening:
  - can create or speed up convergence
  - avoids exploding or going in circles around the solution
  - See Figure 3.4 in the textbook for an illustration.
- $\omega > 1$ : Extrapolation – speeds up convergence (Fig. 3.5)

## Figure 3.4

Dampening to Stabilize an Unstable “Hog Cycle”.


- Suppose inverse demand is  $p = 21 - 3q$  and supply is  $q = p/2 - 3$
- Linear system is not diagonally dominant:

$$\begin{pmatrix} 1 & 3 \\ 1 & -2 \end{pmatrix} \begin{pmatrix} p \\ q \end{pmatrix} = \begin{pmatrix} 21 \\ 6 \end{pmatrix} \quad (3.9.8)$$

- Gauss-Seidel is unstable:

$$p_{n+1} = 21 - 3q_n \quad (3.9.9a)$$

$$q_{n+1} = \frac{1}{2}p_{n+1} - 3 \quad (3.9.9b)$$



# Figure 3.5

## Exatrapolation to Accelerate Convergence in a Game

- Assume firm two's reaction curve is  $p_2 = 2 + 0.80p_1 \equiv R_2(p_1)$ , and firm one's reaction curve is  $p_1 = 1 + 0.75p_2 \equiv R_1(p_2)$ .
- Equilibrium system is diagonally dominant
- Gauss-Seidel is the iterative scheme


$$p_1^{n+1} = R_1(p_2^n) \quad (3.9.12a)$$

$$p_2^{n+1} = R_2(p_1^{n+1}) \quad (3.9.12b)$$

- Accelerate (3.9.12). If  $\omega = 1.5$ , we arrive at faster scheme:

$$p_1^{n+1} = 1.5R_1(p_2^n) - 0.5p_1^n, \quad (3.9.13a)$$

$$p_2^{n+1} = 1.5R_2(p_1^{n+1}) - 0.5p_2^n. \quad (3.9.13b)$$



# Dampening and convergence

- If all eigenvalues  $\sigma(G)$  are real, optimal dampening is

$$\omega^* = \frac{2}{2 - \max |\sigma(G)| - \min |\sigma(G)|}$$

- Weyl's inequality states that
$$\sigma\{\omega G + (1 - \omega)I\} \leq \omega\sigma(G) + (1 - \omega)$$
  - Dampening shifts (potentially complex) eigenvalues towards 1
  - If  $\max \operatorname{real}\{\sigma(G)\} < 1$ , dampening can restore stability and improve convergence speed.
  - If  $\max \operatorname{real}\{\sigma(G)\} > 1$ , dampening cannot restore stability.
- Can combine dampening with G-S (Successive overrelaxation)
  - Potentially even bigger gains

# Preview: nonlinear GS/GJ methods

- We can (and will) use iterative methods for nonlinear fixed-point problems:  $x = F(x)$
- For  $\sigma(G)$  and  $\rho(G)$ ,  $G$  is the Jacobian of  $F$
- Convergence properties are local, and we can have more than one solution
- Dampening works exactly as described
- In general, convergence does not necessarily imply a solution
- For linear problem, we can **verify** solution by computing

$$\|Ax - b\|$$

- For nonlinear fixed point problem, no general test

# Alternatives: Krylov subspace methods

- eg *Conjugate Gradient* for pos. def.  $A$ : *GMRES* for general  $A$
- Combines strengths of iterative and direct methods
- "Exact" convergence in  $n$  iterations by building solution from sequence of linearly independent vectors
  - Floating point error accumulates, worse than direct methods
  - Practical implementations require fixes: see Golub & van Loan
- Each iteration uses  $A$  only in small # of matrix-vector multiplies
- No real gain over direct methods if no structure:  $n \times O(n^2)$  ops
- Substantial gain in speed if
  - 1 Matrix-vector multiplies are fast: sparse, structured, etc
  - 2  $b$  (nearly) restricted to low-d subspace, or initial guess good
  - 3 Preconditioning can help ensure this
    - Cut off at small # of iterations for small loss in accuracy



# Parallelization

- GJ with single CPU computes  $x_i^{k+1}$ 's one after another
- If we have many CPUs, we can parallelize:
  - CPU1 computes  $x_1^{k+1}$ ,
  - CPU2 computes  $x_2^{k+1}$ , and so on
  - These computations happen simultaneously
- Matlab has parallelization built-in, but it needs to be activated
  - In other languages, look for MPI libraries
- Not very useful with just 2 CPUs

# Randomized Methods

- Random sampling of data points or random feature transforms
- $P$   $m \times n$  random projection or transform,  $m \ll p$
- Replace  $A$  by  $P * A$  or  $A * P'$  or even  $PAP'$
- Reduce effective size of data matrices in econometrics/ML applications
- Trade off probability of error for gains in speed and memory
- Especially attractive if approximate structure exists
  - Sparse, low rank, or i.i.d. matrices amenable to accurate random transforms

# Structured linear problems – Sparse methods

- Matrix is "sparse" if most elements are zero
  - Can store only nonzeros elements (and their  $i, j$ )
- Skipping zero elements saves computation time
  - Computing  $0 * 0$  would take same time as  $3.14 * 2.72$
- To take advantage of it, software must be able to:
  - Store matrix as a list of nonzero elements
  - Use only nonzero elements in linear algebra operations
- *C, Fortran, Julia, Python* have sparse linear algebra libraries
  - Julia: `Pkg SparseArrays`, Python: `Scipy.sparse`
- Matlab/Julia have integrated support for sparse matrices:
  - `As=sparse(A)`;  $\Rightarrow$  most operations on `As` will use sparsity

# Sparse example – Markov transition

- Economic system with discrete states:  $i = 1, \dots, 100$ :
  - Markov process: state transition depends only on current state
  - In each period, system can transit to a neighboring state, or stay in current state
- Markov transition matrix
  - $\pi_{ij} = \Pr\{\text{state} = j \text{ tomorrow} \mid \text{state} = i \text{ today}\}$
  - $\Pi = \{\pi_{i,j}\}$  has only 3 non-zero's per row, i.e. 97% of zeros:
  - If  $x$  is the current distribution over states , then in the next period, we will have distribution  $\Pi x$
- Computing  $\Pi x$ :
  - Normally requires  $100^2$  multiplications & additions
  - Accounting for sparsity, we have only 300 multiplications
- Another example: stationary distribution solves  $\Pi x = x$ .

# Eigenvalue/vector computation

- Stationary distribution is example of an *eigenproblem*
- Find vector and complex scalar  $(v, \lambda)$  s.t.  $Av = \lambda v$
- Applications in dynamical systems & control (stability), statistics (PCA, CCA, etc), portfolio management
- Core component of several other problems: eg SVD
- Direct and iterative methods available here too
- Direct methods apply matrix decomposition
- Schur:  $A = UTU^*$ ,  $T$  upper triangular,  $U$  unitary
  - Eigenvalues read off of diagonal of  $T$
  - Computable in  $O(n^3)$  by repeated  $QR$  decompositions
  - Eigenvectors require additional  $O(n^3)$  step: avoid if not needed
- Matlab/Julia/Python: use `eig()`/`eigvals()` for all eigenvalues by direct method

# Iterative eigenvalue computation

- Especially useful if not all eigenvalues needed
- Can find one at a time: often just need biggest or smallest
- Simplest method: *Power iteration*
  - $z_{k+1} = A * z_k / \|A * z_k\|$
  - Converges to eigenvector of largest eigenvalue if that is unique and if  $z_0$  not orthogonal to it
  - Interpretation: Evolution of distribution under markov chain
  - Since uses multiplications, fast if sparse or structured
- Generalized power methods (Arnoldi or Lanczos, etc) improve speed especially for multiple eigenvalues
- Matlab: `eigs()` for subset of eigenvalues by iterative method
- Julia: `eigs()` in Arpack, or option in `eigvals()` for range

# Recommended References: General Numerical Linear Algebra

- Judd Chapter 3.
- Quantecon "Numerical Linear Algebra" [https://julia.quantecon.org/tools\\_and\\_techniques/numerical\\_linear\\_algebra.html](https://julia.quantecon.org/tools_and_techniques/numerical_linear_algebra.html)
- Golub, Gene and Charles van Loan. Matrix Computations
  - The source for matrix algorithms: LAPACK based on them
- Carron, Igor. *The Advanced Matrix Factorization Jungle* <https://sites.google.com/site/igorcarron2/matrixfactorizations>
- Higham, Nick. *Blog* <https://nhigham.com/blog/>
  - Linear algebra facts and algorithm explanations

# Specialized Topics References

- Iterative Methods
  - QuantEcon "Krylov Methods and Matrix Conditioning"  
[https://julia.quantecon.org/tools\\_and\\_techniques/iterative\\_methods\\_sparsity.html](https://julia.quantecon.org/tools_and_techniques/iterative_methods_sparsity.html)
  - Goh, Gabriel. "Why Momentum Really Works" *Distill* 2017  
<http://distill.pub/2017/momentum>
- Randomized methods
  - Mahoney, Michael. "Randomized algorithms for matrices and data" *Foundations and Trends in Machine Learning: Vol. 3: No. 2*, pp 123-224 <https://arxiv.org/abs/1104.5557>
  - Ng, Serena. "Opportunities and Challenges: Lessons from Analyzing Terrabytes of Scanner Data" 2017 NBER WP23673  
<http://www.nber.org/papers/w23673>